

Competing Explanations, Persisting Rules: Grounded Defense against Misleading Charts in Executing Web Agents

Anonymous submission
Paper ID: NAACL-2027-XXXX

Abstract

Vision-language agents that execute multi-step web tasks treat charts as evidence: the way a chart is read enters tool calls, submissions, and actions. Charts, however, admit readings that are both faithful to the rendered image and wrong about the underlying data, because correct visual observation and misleading encoding conventions coexist in the same figure. Existing defenses operate at the level of single question-answer rounds and do not address how an interpretation, once revised, remains in force over a stateful execution. We present CERS, a harness-level defense that organizes chart interpretation around two coupled mechanisms: a competing-explanation chain that expresses the current judgment as $O \wedge B \Rightarrow C_q(a)$ —an observation, an interpretation rule, and the conclusion an action requires—and a persistent rule state that records every validated rule revision with its supporting evidence and scope, and re-uses it in later reasoning. On a paired static panel of 140 executable web tasks built on MISVISAGENT-BENCH, CERS improves the robust accuracy of a 27B open-weight VLM agent from 57.9% to 68.6% in the original round and replicates the gain in an independent round (52.1% to 64.3%; McNemar $p = 0.0025$), with a smaller gain on an 8B model (46.4% to 57.1% for the most compatible variant). In stateful online execution over a closed deployment of twelve live tasks spanning both chart variants, both systems submit only correct actions when they submit at all (26/26 and 27/27 in two frontier-model batches); the defense arm converts ordinary agents’ misleading-chart decisions into recorded abstentions rather than wrong submissions, and a stability rerun agrees with the formal run on 37 of 48 unit outcomes, with every disagreement an abstention–correct flip.

1 Introduction

Large language model agents act for their users by looping through observation, reasoning, and action

in an external environment. What such an agent believes about the information it reads enters tool calls and operations directly, which makes the reliable interpretation of external content a precondition for execution safety. Prior work has shown that indirect prompt injections buried in web pages and tool returns can steer an agent away from the user’s intent (Zhan et al., 2024; Evtimov et al., 2025; Johnson et al., 2025), that perturbations of visual input can hijack multi-step web agents (Wu et al., 2024), and that deceptive interface design can be reproduced and probed systematically (Wang et al., 2025). These findings indicate that an environment does not only intervene on an agent through explicit instructions; it can also act on the way the agent understands and uses information, and thereby shape the decisions the agent makes autonomously.

Charts are a prominent part of such external information. They organize data into visual forms that support comparison and inference, and they lend an appearance of evidential support to whatever judgment is attached to them. Yet the basis for a correct reading and a misleading presentation can coexist in the same figure: research in visualization shows that charts are used to support misleading inferences even when the data are drawn faithfully (Lisnic et al., 2023; Nadib et al., 2026). A curve that moves upward across the rendered image is observed accurately, for instance, but whether it represents growth depends on the axis mapping currently in force; the meaning of color intensity and bar height is likewise read through the encoding relations of the present figure. A misleading presentation can induce an agent to apply its customary reading rules unchanged, converting genuinely visible evidence into a wrong conclusion. Chen et al. (2025) document that multimodal models are systematically vulnerable to such manipulations in chart question answering, (Mahbub et al., 2025, 2026) trace the vulnerability and its partial mitigation across model families, and our companion

work under review formalizes the execution-level consequence: on a paired benchmark of executable web tasks, biases induced by misleading charts propagate into subsequent actions and change task outcomes.¹

Against misleading charts, model-level defenses have improved chart question answering through tabular answering and re-plotting procedures that force a structured restatement of the evidence (Ton-glet et al., 2026). These procedures provide an effective basis for model-level repair, but the reliability of the verification and correction itself remains open. In the inference-only setting, intrinsic self-correction can damage answers that were initially correct (Huang et al., 2023), and when the evaluation component is itself exposed to the environment, reflection and search can introduce new attack surfaces rather than remove existing ones (Wu et al., 2024). Chart-domain studies observe a further gap: a model may correctly identify chart properties during extraction and still fail to let those properties govern the downstream inference (Mah-bub et al., 2026). These observations point to two connected requirements for a defense: correction must rest on evidence that can be checked, and a corrected interpretation must actually change subsequent judgments. Repeated observation or an additional review round can discover an error, but it can equally re-adopt an inapplicable reading, or overwrite a judgment that was already correct.

Stateful execution raises requirements beyond chart question answering: an understanding formed from one chart feeds later planning and operations even after the chart has left the viewport. Once an axis mapping has been verified, later judgments should keep using that verification; when sufficient new evidence appears, the same interpretation should become revisable. Reflexion-style memory carries verbalized feedback into later attempts (Shinn et al., 2023), and recent agent-memory systems organize experience into retrievable state (Wang et al., 2024; Hu et al., 2025); an inductive chart requires something more specific, the current validity of interpretations: which rules have evidential support, which have been refuted, and what new evidence suffices to change an accepted reading? We pose one unified question: *how can an agent reliably revise visual interpretation rules,*

and keep validated rules in force over subsequent execution?

We answer this question with a defense that combines *competing explanation chains* with a *persistent rule state*. When the agent forms a chart-dependent judgment, the harness expresses it as $O \wedge B \Rightarrow C_q(a)$: an observation O in the chart, an interpretation rule B that turns the observation into task meaning, and the conclusion $C_q(a)$ that an action a requires under the task q . The harness then builds candidate chains with different interpretive bases on the same complete chart, turning default visual conventions into explicit objects of test, and compares their applicability using evidence drawn from the figure itself. Validated rule revisions enter the agent’s runtime harness as an interpretation state that later reasoning reads and updates by scope; subsequent decisions reuse the rules currently in force, and re-enter verification only when new effective evidence challenges them. Competing explanations determine why an interpretation needs to change; the persistent state carries that evidenced change into later actions.

The contributions of this paper are two-fold. First, a correction paradigm based on competing explanation chains: the relation $O \wedge B \Rightarrow C_q(a)$ unifies visual observation, interpretation rules, and action judgments; competing chains are constructed on one chart, and the applicability of customary readings is tested against in-chart evidence, so that the rules responsible for wrong inferences are identified and revised. Second, a stateful defense for continuing execution: validated rules, their revision evidence, and their scope conditions are maintained as task-scoped state inside the harness, so that later decisions reuse confirmed chart readings and revise them only on new effective evidence.

2 Related Work

Research on misleading visualizations has developed along two lines that our defense connects. The first line studies how charts mislead human readers: Lisnic et al. (2023) show that deception in charts relies less on outright visual distortion than on the combination of faithful drawings with selective context, framing, and commentary, and current visualization research extends this analysis to persuasive chart design (Nadib et al., 2026). The second line measures the same phenomenon in models: benchmarks evaluate multimodal LLMs on misleading chart question answering (Chen et al.,

¹Companion submission under review; the paired benchmark MISVISAGENTBENCH and the execution-level evaluation protocol it introduces are described where first used in this paper.

2025; Bharti et al., 2024; Alexander et al., 2024), Mahbub et al. (2025) and Mahbub et al. (2026) quantify how truncated axes, distorted aspect ratios, and manipulated encodings degrade chart reasoning across model families and probe mitigations, and recent studies examine whether agents and LLMs recognize misleadingness without being nudged (Panda, 2026; Blasilli and Angelini, 2026; Seto et al., 2026). Method-side responses include tool-augmented chart reasoning that grounds answers in extracted chart elements (Kaur et al., 2026), structured extraction pipelines that repair extraction–inference gaps (Das et al., 2025), detection of misleading charts as a classification task (Tonglet et al., 2025), and agentic dual-path frameworks that separate data fidelity checking from conclusion checking (Zhang et al., 2026). Tonglet et al. (2026) provide the closest model-level defense, protecting multimodal LLMs through tabular answering and chart re-plotting; these procedures repair a single question–answer round, while CERS targets the interpretation rules that recur across the rounds of an executing agent and maintains their validated revisions in runtime state.

A second related line concerns the integrity of agent environments. Indirect prompt injection benchmarks show that content returned from tools and web pages can steer tool-augmented agents (Zhan et al., 2024), web-agent security benchmarks measure targeted hijacking under realistic threat models (Evtimov et al., 2025), and demonstrations show injection through the very representations agents perceive, such as HTML accessibility trees (Johnson et al., 2025); automated red-teaming generates such attacks at scale (Wang et al., 2025) and runtime frameworks intercept them (Zhao et al., 2026). Defenses at the interface level separate instructions from data through structured queries (Chen et al., 2024a) or preference-optimized training against injected content (Chen et al., 2024b). Wu et al. (2024) dissect the robustness of multimodal agents as compound systems and show that reflection and tree search, the components commonly added for self-correction, open new vulnerabilities when the evaluator itself is exposed to attacked content. Our defense differs in what it protects: not the agent’s instruction hierarchy but the interpretation layer between visual observation and task judgment, where a misleading chart exerts its effect without any injected instruction.

The third line concerns persistent state and correction in agents. Verbal reinforcement learning

carries textual feedback across attempts (Shinn et al., 2023), self-consistency aggregates multiple reasoning paths into one answer (Wang et al., 2022), and intrinsic self-correction without external feedback degrades reasoning performance (Huang et al., 2023). Agent surveys organize the growing design space of state and memory components (Wang et al., 2024; Hu et al., 2025), and Koh et al. (2024) establish that realistic visual web tasks remain hard enough that correctness of single decisions cannot be taken for granted. CERS borrows the state-carrying idea from memory research but changes the carried content: instead of outcome feedback or free-form reflections, the state records interpretation rules with their evidential support, applicability scope, and revision history, so that later steps inherit what was verified about the chart rather than a verdict about the answer.

3 Method

CERS rests on one object: the interpretation rule B that connects visual observation to task judgment. Competing explanation chains make these rules explicit and test them; persistent state organizes their revision evidence, current validity, and applicability scope, so that they participate in later interaction. This section develops the two mechanisms in turn after fixing the task setting.

3.1 Task Modeling and Overall Framework

Consider an agent executing a user task q . At step t the agent receives an environment observation o_t and, given the interaction history H_t , proposes an action $\tilde{a}_t$. When the proposal depends on chart understanding, let I_t denote the complete chart observation relevant to the current judgment; if the chart has left the current page, I_t refers to the chart observation actually obtained earlier in the task that still bears on the current judgment. We express a chart-driven interpretation as

$$O \wedge B \Rightarrow C_q(a), \quad (1)$$

where O is an observation in the chart, B is the interpretation rule required to convert the observation into task semantics, and $C_q(a)$ is the proposition that the action, or high-level choice, a meets the requirements of task q . The representation separates a visible phenomenon from the condition under which it is interpretable: the observation that a curve moves upward within an interval, jointly with the rule that in the current axis a higher position means a larger value, supports the conclusion

that the interval grows and that the corresponding action is admissible. For a geometric distortion, the observation may be that object A’s bar stands higher than object B’s, while the interpretation rule is that the current bar-height ordering preserves the value ordering; if the verified numeric relations in the figure disagree with the geometric ordering, it is precisely the applicability of this rule that is challenged first. Both O and B may carry multiple items; they are expressed in structured natural language covering entity binding, numeric relations, encoding direction, time range, and comparison conditions. Relations that can be computed explicitly, such as numeric comparisons and orderings, are checked with program assistance; the remaining semantic judgments are verified jointly with the public task requirements.

The method maintains a within-task rule state M_t^B and, for chart-dependent judgments, runs the following process:

current interpretation $\rightarrow$ competing-explanation generation $\rightarrow$ in-chart applicability verification $\rightarrow$ rule-state update $\rightarrow$ ac
(2) $\Pi_t = \{\pi_0, \pi_1, \dots, \pi_K\}, \quad \pi_i = \langle O_i, B_i, p_i \rangle, \quad (4)$

Later steps read the applicable rules from the current context and restart competing verification only when a relevant judgment is first formed, when new evidence conflicts with an existing rule, or when the applicability conditions of a rule are in question. The process involves two intertwined time scales. Local verification decides whether a particular interpretation rule deserves acceptance or revision; within-task maintenance decides in which later steps that rule remains effective. The parameters of the underlying model stay frozen throughout, and the rule state evolves with the observations and verification results actually obtained.

3.2 Competing Explanation Chains

3.2.1 Making the Current Explanation Explicit

When the agent proposes a chart-dependent choice $\tilde{a}_t$, the interpretation module forms, from the complete chart, the public task, the admissible history, and the current rule state,

$$\pi_0 = \langle O_0, B_0, p_0 \rangle, \quad p_0 = C_q(\tilde{a}_t), \quad (3)$$

stating separately which phenomena the chart is claimed to contain, which rules carry the observation to the conclusion, and which evidence those rules already hold or lack. Where the current choice lacks support, the module may record an interpretation gap directly. A statement such as “the color

is darker” can stand as an observation, whereas “the risk is higher” stands only if a mapping from color to risk holds; making the chain explicit turns that omitted inference condition into an object for later verification. The chain is a checkable record of the task argument: it exposes which conditions the current choice requires and directs the counter-questions that follow.

3.2.2 Counter-Questions and Competing Explanations

The competing module poses three connected questions around the current chain: whether the current observation suffices for the conclusion, which interpretation rules the inference depends on, and whether the same chart admits another evidence-based explanation under which the current conclusion fails, or under which the current action continues to hold on different grounds. From these questions the module generates a bounded candidate set

where each conclusion p_i may support the current action, refute it, or support a different task judgment. All candidates refer to the same complete chart, and the counter-questions act on the interpretive conditions between observation and conclusion. A competing chain is meaningful only if it differs substantially in observation correspondence, interpretation rule, applicability condition, or scope of derivation: two chains that both accept “the curve rises” but adopt forward and reversed value mappings disagree precisely on the applicability of B . A reversed interpretation enters first as a conditional hypothesis—if the position–value relation of the current axis differs from the customary direction, the original trend judgment must change—and remains unverified until the chart’s scales and correspondences support the condition. The competing module may return insufficient support and is not required to produce an opposing conclusion for every task.

3.2.3 Normalization and Shared Rules

Generated chains are normalized over observation, rule, applicability scope, conclusion, and evidence dependence. Explanations that differ only in wording are merged, and a rule used in several chains shares one verification state; different entities or comparison relations keep their separate evidence. “Darker means higher risk” and “a deeper color

stands for a more severe risk” map to one rule when object and scope coincide, whereas comparing Illinois with Kansas and Illinois with Delaware under the same legend are two comparisons that shared rules do not merge. Whereas self-consistency aggregates a distribution of answers over repeated reasoning paths (Wang et al., 2022), the object of verification here is the rule each path depends on and the evidence for it: repeated generations of the same rule add no evidential weight, and a new explanation changes later verification only by contributing different evidence or exposing a new applicability problem.

3.3 In-Chart Applicability Verification

The verification module first extracts the disputes that genuinely affect conclusions among the candidate chains. If several chains agree on the observation and disagree only on mapping direction, verification concentrates on that mapping; if the dispute concerns entity–value correspondence, observation binding is checked before rule applicability is judged. For each chain the module outputs

$$V_i = (v_i^O, v_i^B, v_i^{\Rightarrow}, E_i), \quad (5)$$

where each verification status takes the values support, refute, or undetermined, and E_i records the in-chart evidence. Observation O_i is checked for presence in the chart and for the correctness of entity, value, time, and unit binding; rule B_i is checked for support by the current chart and satisfaction of its applicability conditions; the derivation $O_i \wedge B_i \Rightarrow p_i$ is checked for sufficiency of the confirmed premises, including candidate coverage, range, and task conditions. Verification always proceeds against the complete chart, and the evidence record contains the actual observation identifier, the content read, and which rule the content supports or refutes. To keep the judgment focused on evidence, verification takes the normalized dispute propositions as input, and the source role or repetition count of a candidate is never used as a credibility signal; the module confirms observations and interpretive conditions before evaluating conclusions, so that a preferred answer is not used to justify a rule retroactively.

When a chain is challenged, the method distinguishes three objects of update. An *observation error*—a misread number, a wrong entity binding, a wrong range—revises the affected O and re-verifies the interpretations that depend on it; unreliable ob-

servations cannot be charged to B directly. An *inapplicable rule* is one whose observation stands while the mapping from observation to semantics disagrees with the current chart; such a rule is revoked, or its applicability narrowed, within the current scope. *Insufficient derivation* leaves both observation and rule supported while the full task conclusion remains unproven; the module then adds the missing comparison or keeps an undetermined status, as when one higher candidate refutes the claim that the current candidate is highest without yet confirming that the alternative is globally highest. This separation keeps the defense centered on interpretation rules while preserving correct handling of perception, scope, and task-derivation problems.

Conflicting conclusions between candidate chains establish only that a conflict exists. A chain is excluded only when chart-supported evidence negates its observation, its rule applicability, or its derivation conditions. For rules that are mutually incompatible over the same objects and scope, the method checks which in-chart information can discriminate them: if two explanations each retain unresolved evidence, the dispute is kept open; if they apply to different axes, series, or ranges, the applicability conditions are split rather than forced to a single choice. Verification may support a replacement rule, or only negate the original one: discovering that bar height and visible value ordering disagree does not entail that “shorter means larger,” and the sound revision may instead be that bar height in the current figure does not directly support quantity ordering, which must rely on verified entity–value relations. Rule revision therefore includes both replacing a mapping and disqualifying an inference route within the current chart.

3.4 Rule Revision and Persistent State Update

Each interpretation rule is maintained as a record

$$m_i = \langle \text{id}_i, B_i, \sigma_i, s_i, E_i, \Delta_i, v_i \rangle, \quad (6)$$

where id_i identifies the rule for later reference, B_i is the current interpretation rule, σ_i is its applicability scope over chart objects, encoding components, and semantic conditions, s_i is its status among pending, currently active, disputed, and revoked, E_i is the observed evidence supporting or refuting the rule, Δ_i links the original rule, the revised rule, and the reason for the change, and v_i is the revision version. The persistent state stores how the current chart is to be read and on what evidence;

concrete object choices, task answers, and click actions are produced by the task reasoning itself and recorded in the execution history. A rule’s scope is fixed by the conditions under which it holds: a position–value mapping binds to the numeric axis and series concerned, a color reading binds to the current legend and metric, and a rule about overall trend may further involve temporal order and comparison range. Page changes unrelated to the task goal do not invalidate a rule by themselves.

Verification produces rule-revision proposals Δ_t^B , which the state manager applies as $M_{t+1}^B = U(M_t^B, \Delta_t^B)$. When the original rule is supported and the new proposal lacks effective evidence, the original rule is kept and the proposal is retained as a pending record. When the original rule is refuted and a replacement is supported, the original loses applicability in the affected scope and the replacement version is activated. When the original is refuted without a settled replacement, the rule is revoked and the current interpretation gap is recorded. When the original and a competing rule remain in effective conflict, the affected rules enter a disputed status and verification continues within budget, and when two rules apply to different objects or scopes they are stored separately rather than overwritten globally. The “currently active” status expresses the method’s judgment under the evidence at hand; new effective evidence can change it, and can also show that an earlier verification was itself mistaken. After revocation, a rule’s content, scope, and reason remain in the revision history; if the same rule reappears later, it must answer the recorded refutation or supply evidence that changes the applicability judgment, since repetition, later position, or a more confident phrasing forms no new revision ground by itself.

A revision record answers why the original rule had to change, what interpretation is adopted now, and under which conditions the change applies: for a reversed linear axis, the record replaces “higher position, larger value” because the visible scale increases downward, states that an upward move of the bound series now reads as a decrease, and scopes this to the current chart, axis, and series—supporting trend judgments over further intervals without storing a fixed answer of growth or decline.

An invalidated rule does not by itself condemn the current action, since $\neg B_0 \not\Rightarrow \neg C_q(a_0)$. After updating B , the method re-examines the current judgments that depended on it: a judgment that lost its basis enters re-review and may be retained

if other effective interpretations still support it, so the two admissible outcomes are a rule change with the action retained (a wrong reason for a right choice) or rule and action changing together. Judgments already recorded in the history are marked as basis-changed for later re-confirmation; the original interaction history stays intact, and the currently effective rules are organized by a separate read interface.

3.5 Reading and Using Rules over Continued Execution

Let the current interpretation context c_t collect the relevant chart objects, encoding components, metrics, and task scope. The state reader returns $\mathcal{R}_t = \text{Read}(M_t^B, c_t)$, of which the directly usable part is

$$\mathcal{R}_t^{\text{active}} = \{m_i \in M_t^B : s_i = \text{active} \wedge \text{Match}(\sigma_i, c_t) = \text{true}\}. \quad (7)$$

Scope matching yields applicability, non-applicability, or undetermined; an undetermined match triggers verification of the corresponding condition rather than a default assumption that a stored rule covers the new chart. The read result also carries the disputes and revocation records relevant to the current judgment, so the agent can separate interpretations still in force from ones that appeared before and lost support; revoked content never participates in derivation as a current fact. When a chart has temporarily left the viewport, the reader keeps supplying applicable rules through the chart references established within the task, and new chart objects or encoding changes are bound and verified afresh.

The base agent produces its next proposal from the current observation, history, and rule view,

$$(\tilde{a}_t, \tilde{\pi}_t) = \mathcal{F}_\theta(q, o_t, H_t, \mathcal{R}_t), \quad (8)$$

where $\tilde{\pi}_t$ is a short interpretation record of the present judgment. Rules about the same chart can remain valid while the correct action changes with task conditions, so the method reuses interpretation evidence across steps while re-deriving actions from the current observation and the public task. If a new proposal explicitly depends on a rule that is revoked within its scope, the harness returns the relevant revision evidence and requires re-derivation; if the proposal brings new evidence sufficient to challenge an active rule, competing verification is entered again. This treatment both restrains old conventions from regaining force without grounds

and preserves a channel for correcting earlier verification mistakes. Ordinary navigation, input, and other actions that do not involve a change of chart interpretation follow the original execution interface, and the persistent state acts mainly through later contextual reads, with complete competing verification executed only when a trigger condition holds.

Three observable entry points start verification. The first chart-dependent judgment of a task expands the current interpretation into $O \wedge B \Rightarrow C_q(a)$ and establishes the task’s rule state. An interpretation conflict re-verifies the disputed part whenever a later proposal uses a rule inconsistent with the effective state, or an existing rule is challenged with in-chart evidence. An applicability change checks whether existing rules still hold when the chart, its encoding, the series binding, or the task scope changes in ways that can affect interpretation. The entry points are determined by actual observations, task context, and action proposals; task correctness and hidden scoring never participate in triggering.

After a rule update, the agent re-judges the action under the currently effective interpretation and records one of three outcomes. *Keep* means the current choice continues to be supported by an effective interpretation and the task conditions. *Revise* means the current choice is negated by relevant evidence and an alternative choice gains the support the task requires. *Unresolved* means an interpretive dispute that could still change the action judgment persists, or necessary evidence is missing. Showing that some candidate exceeds the current choice refutes the claim that the current choice is highest, for example, while confirming the replacement still requires the comparison that the public task demands; the method treats such checks as part of the task-semantic verification rather than imposing one fixed comparison formula on all tasks. The unresolved status continues whole-chart verification within the remaining budget, and when the budget is exhausted the unfinished reason is recorded explicitly. Verification suggestions, actual choice revisions, and final submissions are recorded separately, and only operations carried out by the executor constitute action outcomes.

3.6 Full Algorithm and Integration

Algorithm 1 (Appendix A) summarizes one execution of CERS under the candidate budget K and the review budget L for a single verification node. The framework enters an existing agent at three

points: the observation entrance establishes chart references, the pre-execution point performs verification and rule revision, and the later-reasoning entrance reads the currently applicable rules; planning, execution, and review modules share the rule state within one task, while action capabilities and the environment interface follow the original system. Explanation generation and visual verification can be carried out by different calls of the same frozen model. Deterministic code handles structured record keeping, management of normalized merges, version updates, scope retrieval, and checks of relations that can be computed explicitly. The competing-chain count, review rounds, and trigger settings are fixed before a run, and the extra calls and context overhead are charged to the method’s cost. The rule state is isolated per task: the runtime uses only the chart and interaction information the task permits, while the paired chart, hidden answers, and scoring feedback stay on the evaluation side. Charts, rules, and action records are stored separately so that later analysis can trace how evidence changed and whether operations followed.

4 Experiments

This section reports the two evaluation deployments of CERS: a static paired panel that isolates the verification core, and a stateful online deployment that adds the persistent rule state inside a live browser harness.

4.1 Setting

The task basis is inherited from our companion benchmark MISVISAGENTBENCH, which pairs 140 executable web tasks—47 business, 35 environment, 19 health, 39 public-statistics, in 21 mechanism families—with a faithful chart and a misleading chart whose presentation violates an encoding convention while the underlying data remain drawable.² Model inputs never contain gold labels, mechanism labels, the other chart variant, or scoring feedback.

Static panel. For every task the agent sees the original chart of the benchmark—the misleading instance of the pair, scored against its frozen original label—plus the public goal and four options, and must submit a selection; no browser navigation occurs. Six configurations are compared on iden-

²Companion submission under review; instance construction, taxonomy, and scoring are defined there.

tical inputs and decoding (temperature 0.7, top- p 0.8, top- k 20, fixed seed, thinking disabled): direct selection and five verification configurations implementing the READ–VERIFY–DECIDE pipeline of Section 3 with frozen prompts, three requests per task, and a fixed pool of 275 pre-generated observation/rule candidates. They differ in verification emphasis: V3 verifies O and B in distinct schema-typed passes; V4 fixes explicit numeric references, dates, ranges, and baselines; V5 (full) constrains query points with same-series references, checks scales, time windows, and policy boundaries, and prioritizes the discriminating scope; V6 checks each necessary condition of B item by item, requiring positive evidence for routine defaults; V7 keeps V4’s prompts and reorients the decision phase toward same-series position and numeric ranges to avoid abstention on missing exact printed values. Interface failures and empty selections remain in the 140 denominator.

Online stateful deployment. Twelve tasks spanning all four domains (six business, three environment, two health, one public-statistics) run as live browser sessions under a closed deployment: for each task, both chart variants (official misleading chart, clean chart) execute under both systems—ordinary execution without defense, and the CERS arm whose hooks, rule state, and re-derivation loops run inside the harness with the actor driving a real browser. The submission interface is a strict typed protocol (interface `typed_scalars_v2`, at most eight actor calls, one re-verification round, temperature 0), in which a unit ends as a submitted action scored by the original evaluation, an abstention with a recorded reason, or a failure. Three actor systems are used: a local 27B open-weight VLM served with strict response-format constraints (Batch A), the same model behind a schema-guided proxy that mirrors the sealed structural rules of the protocol (Batch B), and the frontier model of the pilot accessed directly through the same gateway (Batch C, the formal result, plus one stability rerun under identical configuration).

4.2 Static Panel Results

Table 1 reports the original round under the 27B open-weight VLM: direct selection sits at 57.9% (81/140), and the full configuration V5 reaches 68.6% (96/140), a net gain of 15 decision points (10.7 points of accuracy) with no increase in abstention or interface failure. Because the panel is entirely single-decision-point, that gain cannot

Config	Target	Rate	Other	Empty	IF
Direct	81	57.9%	59	0	0
V3	74	52.9%	60	0	6
V4	87	62.1%	53	0	0
V5 (full)	96	68.6%	44	0	0
V6	84	60.0%	56	0	0
V7	91	65.0%	49	0	0

Table 1: Original round on the 27B open-weight VLM (140 tasks; temperature 0.7, top- p 0.8, top- k 20, fixed seed). “Target” counts selections matching the original label; “Other” counts wrong selections; “IF” counts interface failures, which remain in the denominator. The 24-task development subset is scored separately and reported in Appendix B.

Config	27B replication			8B cross-model		
	Tar.	Δ	p	Tar.	Δ	p
Direct	73	—	—	65	—	—
V3	75	+2	.436	50	−15	.999
V4	86	+13	.0049	77	+12	.158
V5	90	+17	.0025	75	+10	.089
V6	86	+13	.0014	71	+6	.465
V7	89	+16	.0049	80	+15	.043

Table 2: Replication on the same 27B model (left; 8.84M tokens, 2,230 requests) and transfer to an 8B open-weight VLM (right; 8.14M tokens, 2,184 requests). Δ is the paired net change against the Direct configuration of the same run; p values are two-sided exact McNemar tests on discordant decisions.

be attributed to information gained after selection: verification changes the decision that was already available, which is the quantity the formulation of Section 3.1 isolates.

The replication run keeps the model, seeds, and prompts frozen and re-executes the whole pipeline (Table 2, left). Direct selection drops from 81 to 73 of 140 (57.9% to 52.1%) under an identical configuration, which quantifies the sampling variance of agentic decoding itself; every comparison must therefore be made within the same run. Within-run, V5 gains +17 paired decisions (McNemar $p = 0.0025$; exact 95% CI on the difference in accuracy $[+0.029, +0.214]$), and V4, V6, and V7 gain +13, +13, and +16 with $p < 0.005$ each, while V3 returns to +2. Paired transfer counts are one-directional in aggregate (57 keep, 33 error-to-target, 11 target-to-error, 5 target-to-empty, 0 empty-to-target), and the panel remains abstention-

free outside V3’s interface failures.

Cross-model transfer. On the 8B open-weight VLM, direct selection drops to 46.4% and the picture splits (Table 2, right). V7 gains +15 ($p = 0.043$; CI [+0.007, +0.207]), V4 and V5 gain +12 and +10 without reaching significance, and V6 +6 is within noise. V3 collapses: 51 of 140 runs end in interface failures of the schema-typed protocol, against 8 in the 27B replication and 6 in the original round, and its net effect is -15 . The interaction is instructive: on the 27B model V3’s separated *O/B* verification produced the largest error-to-target correction count of any configuration (30 against 6 reverse, $p = 0.0001$), for a net of +2 after interface failures; the same prompts overwhelm the 8B actor. Interface robustness is thus a first-class experimental variable, and prompt content and interface contract must be evaluated jointly rather than as a single “method” factor.

Where the gains concentrate. Stratified by mechanism family, V5’s net gains concentrate in families whose misleading power comes from scope and reference manipulation: the subset-extrapolation family (22 tasks) yields +18 in the 27B replication and +11 in the 8B transfer, while the non-linear-axis family is neutral in the replication and -6 in the 8B round, and the dual-axis family -5 and 0 . By domain, business carries the largest share in both rounds (+16 and +4), public statistics is neutral in the replication and positive in the 8B round (0 and +9), health is +6 in both, and the environment domain is the only negative one in both rounds (-5 and -9); the original round shows the same qualitative pattern. The development subset contributes only one decision point to each headline gain (+17 comprises 1 development and 16 remaining tasks; +10 comprises 1 and 9), and all 24 development tasks score above the run’s direct-selection mean in every verification configuration. Costs remain bounded: the replication rounds spend 8.84M total tokens for the 27B round (2,230 requests) and 8.14M for the 8B round (2,184 requests), against 8.86M for the original round, a roughly $2.6\times$ request overhead over direct selection with frozen prompts and no parameter updates.

4.3 Online Stateful Results

A pilot (6 tasks, frontier model, 12 units) already showed the qualitative pattern: the ordinary arm submitted 5 correct selections against 1 for the

Unit class	Cor.	Abst.	Bug.	Fail.
Batch C (formal run)				
Ordinary $\times$ misleading	8	3	–	1
Ordinary $\times$ clean	10	2	–	–
CERS $\times$ misleading	2	7	1	2
CERS $\times$ clean	6	4	–	2
Batch C rerun (stability)				
Ordinary $\times$ misleading	10	1	–	1
Ordinary $\times$ clean	10	2	–	–
CERS $\times$ misleading	2	5	–	5
CERS $\times$ clean	5	4	–	3

Table 3: Online stateful deployment: 12 tasks spanning four domains, each run on the misleading chart and its clean pair under ordinary execution and the CERS arm (48 units per batch). “Cor.” counts submitted actions matching the original label; “Abst.” recorded abstentions; “Bug.” re-verification budget exhaustion; “Fail.” engineering and response-format failures (including actor-call-limit terminations). All submissions across both batches are correct (26/26 and 27/27).

CERS arm. The formal batch (Table 3, top) confirms it at larger scale and adds the important control: the ordinary arm loses 3–5 decision points on misleading charts relative to its own clean-chart performance (8 vs. 10 here; 7 vs. 10 and 7 vs. 12 in the two earlier 27B batches, which additionally contain 3 wrong submissions caused by the misleading chart per batch), whereas the CERS arm submits only correct actions and converts the rest into recorded abstentions—7 misleading-chart abstentions, one re-verification budget exhaustion, and two engineering failures. The paired task-level effect is therefore strictly one-directional in both C batches (ordinary preserves 10 and 13 decisions that the defense arm turns into abstentions; the reverse never occurs), and the aggregate submission precision is 100% for both systems when they do submit.

Interface decomposition. The three actor systems separate structural from semantic failure modes. Under the strict interface without guidance (Batch A), 22 of 24 defense-arm units die in engineering failures, and the ordinary arm terminates in two call-limit exhaustions on misleading charts and two engineering failures on clean charts. Introducing a schema-guided proxy that mirrors the sealed structural rules (Batch B) drives the ordinary actor’s structural failures to zero (12/12 on clean charts) and cuts defense-

arm engineering failures only modestly (21 of 24), revealing that the binding constraint is semantic: exact public-leaf values, bounded summaries, and the question–response count coupling cannot be expressed by any response format. Running the frontier model directly through the gateway (Batch C) removes the structural interface almost entirely—defense-arm engineering failures fall to 4 of 24—and the remaining mass is honest abstention: 12 `unresolved_public_evidence` verdicts, most of them on misleading charts, where verification cannot close the gap between the printed chart and the question’s required scope within the re-verification budget.

Stability. Re-running Batch C under the identical configuration changes 11 of 48 unit outcomes, every one of them a flip between a correct submission and an abstention (agreement 37/48); no wrong submission appears in either run, so the “never submit a wrong answer” property is stable while completion throughput is not. Token budgets are comparable across the two runs (1.80M input and 129K output tokens, 322 attempts and 242 browser operations, versus 1.81M and 115K, 331 and 254), confirming that the stateful machinery operates within a fixed per-unit allowance comparable to the static panel’s per-task cost.

5 Conclusion

We treated misleading charts as a runtime evidence problem for executing web agents rather than a benchmark artifact: the agent’s own reading of a chart becomes part of its state, so a defense must both compete explanations at the moment of reading and keep revised rules in force afterwards. CERS implements this pair of requirements with a competing-explanation chain over visible chart evidence, an in-chart verification pass that constrains query points with same-series references, and a persistent rule state that is revised only through a `KEEP/REVISE/UNRESOLVED` decision with recorded evidence and scope. The static panel shows that this machinery yields a within-run robust-accuracy gain of 10–17 paired decision points on a 27B open-weight agent, replicates under a frozen configuration, and transfers to an 8B agent through the decision-oriented variant, while the failure analysis shows that interface contracts co-determine outcomes as strongly as prompt content. The stateful online deployment shows the complement: under a fail-closed submission proto-

col the defense arm never submits a wrong answer and its residual errors become legible abstentions with recorded reasons, at a throughput cost that our interface decomposition attributes to the semantic contract of the protocol rather than to the reasoning itself. We release the full decision logs, rule states, and re-derivation traces so that future work can measure not only whether agents answer correctly, but which explanation they acted on and whether it survived contact with the chart.

Limitations

This work studies defensive scaffolding around vision–language agents, and its limitations follow from the same design choices that produce the results.

Scale and models. The static panel uses one 27B open-weight VLM and one 8B open-weight VLM under fixed decoding; the online deployment uses the same 27B model under two interface regimes and one frontier model accessed through a gateway. We do not claim that the reported gains generalize to proprietary or future models, to other agent frameworks, or to tasks outside the four covered domains; the 8B transfer already shows that variant choice interacts with actor capability in ways we do not fully predict.

What the labels measure. Robust accuracy is computed against the original option labels of the benchmark, and our agreement checks verify that submitted options and rule states match the recorded logs. Label agreement is not semantic verification: a run whose recorded reasoning matches its outcome may still rest on a misreading that the frozen labels do not separate, and conversely a correct submission may be produced for partially wrong reasons. The rule states we analyze are the agent’s own declared rules, not independently validated facts.

Selection effects and variance. The paired gains are within-run comparisons on a fixed task set; the drop of direct selection from 81 to 73 of 140 under an identical configuration quantifies sampling variance, and we therefore report McNemar tests and exact confidence intervals on paired decisions rather than cross-run rate differences. The 24-task development subset was used to calibrate configuration emphasis and is scored separately throughout; still, all configurations were designed by the authors after inspecting aggregate failure patterns, so the panel should be read as a controlled comparison

rather than a held-out evaluation.

Interface dependence. The online results are conditional on the strict submission protocol. Its semantic contract—exact public-leaf values, bounded summaries, count-coupled responses—cannot be expressed in a response format, and under earlier actor systems the same machinery died in engineering failures rather than abstentions. The deployment therefore demonstrates the behavior of CERS under one operating point (frontier actor, direct gateway, strict typed interface), not a interface-robust defense; the negative environment-domain results of the static panel and the one-directional online paired effect are reported as found.

Runtime footprint. Verification multiplies requests by roughly $2.6\times$ over direct selection and the online arm operates within a fixed per-unit budget (1 re-verification round, 8 actor calls); agents deployed under tighter latency or cost envelopes may not accommodate this overhead.

Reproducibility boundary. The companion benchmark is under review and its instances are not yet public; until release, the results here cannot be independently recomputed from the logs alone, although all decision logs, prompts, and rule states referenced in this paper are preserved verbatim in the submission’s supplementary material.

References

- Jason Alexander, Priyal Nanda, Kai-Cheng Yang, and Ali Sarvghad. 2024. [Can gpt-4 models detect misleading visualizations?](#) In *2024 IEEE Visualization and Visual Analytics (VIS)*.
- Shubham Bharti, Shiyun Cheng, Jihyun Rho, Jianrui Zhang, Mu Cai, Yong Jae Lee, Martina Rau, and Xiaojin Zhu. 2024. [Chartom: A visual theory-of-mind benchmark for llms on misleading charts](#). *Preprint*, arXiv:2408.14419.
- Graziano Blasilli and Marco Angelini. 2026. [True \(vis\) lies: Analyzing how generative ai recognizes intentionality, rhetoric, and misleadingness in visualization lies](#). In *arXiv (Cornell University)*. Cornell University.
- Sizhe Chen, Julien Piet, Chawin Sitawarin, and David Wagner. 2024a. [Strug: Defending against prompt injection with structured queries](#). *Preprint*, arXiv:2402.06363.
- Sizhe Chen, Arman Zharmagambetov, Saeed Mahloujifar, Kamalika Chaudhuri, David Wagner, and Chuan Guo. 2024b. [Secalign: Defending against prompt injection with preference optimization](#).
- Zixin Chen, Sicheng Song, Kashun Shum, Yanna Lin, Rui Sheng, Weiqi Wang, and Huamin Qu. 2025. [Unmasking deceptive visuals: Benchmarking multimodal large language models on misleading chart question answering](#).
- Amit Kumar Das, Mohammad Tarun, and Klaus Mueller. 2025. [Charts-of-thought: Enhancing llm visualization literacy through structured data extraction](#). In *IEEE Transactions on Visualization and Computer Graphics (TVCG), PrePrints 5555*, pp. 1-11, 2025.
- Ivan Evtimov, Arman Zharmagambetov, Aaron Grattafiori, Chuan Guo, and Kamalika Chaudhuri. 2025. [Wasp: Benchmarking web agent security against prompt injection attacks](#). *Preprint*, arXiv:2504.18575.
- Xueyu Hu, Tao Xiong, Biao Yi, Zishu Wei, Ruixuan Xiao, Yurun Chen, Jiasheng Ye, Meiling Tao, Xiangxin Zhou, Ziyu Zhao, Yuhuai Li, Shengze Xu, Shenzhi Wang, Xinchun Xu, Shuofei Qiao, Zhaokai Wang, Kun Kuang, Tiejong Zeng, Liang Wang, and 10 others. 2025. [Os agents: A survey on mllm-based agents for computer, phone and browser use](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7436–7465, Vienna, Austria. Association for Computational Linguistics.
- Jie Huang, Xinyun Chen, Swaroop Mishra, Huaixiu Steven Zheng, Adams Wei Yu, Xinying Song, and Denny Zhou. 2023. [Large language models cannot self-correct reasoning yet](#). *Preprint*, arXiv:2310.01798.
- Sam Johnson, Viet Pham, and Thai Le. 2025. [Manipulating llm web agents with indirect prompt injection attack via html accessibility tree](#).
- Rachneet Kaur, Nishan Srishankar, Zhen Zeng, and Sumitra Ganesh. 2026. [Chartagent: A multimodal agent for visually grounded reasoning in complex chart question answering](#). In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 18488–18555, San Diego, California, United States. Association for Computational Linguistics.
- Jing Yu Koh, Robert Lo, Lawrence Jang, Vikram Duvvur, Ming Lim, Po-Yu Huang, Graham Neubig, Shuyan Zhou, Russ Salakhutdinov, and Daniel Fried. 2024. [Visualwebarena: Evaluating multimodal agents on realistic visual web tasks](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 881–905, Bangkok, Thailand. Association for Computational Linguistics.
- Maxim Lisnic, Cole Polychronis, Alexander Lex, and Marina Kogan. 2023. [Misleading beyond visual tricks: How people actually lie with charts](#). pages 1–21.
- Ridwan Mahbub, Mohammed Saidul Islam, Md Tahmid Rahman Laskar, Mizanur Rahman, Mir Tafseer

- Nayeem, and Enamul Hoque. 2025. [The perils of chart deception: How misleading visualizations affect vision-language models](#). In *2025 IEEE Visualization and Visual Analytics (VIS)*.
- Ridwan Mahbub, Mohammed Saidul Islam, Md Tahmid Rahman Laskar, Mizanur Rahman, Mir Tafseer Nayeem, and Enamul Hoque. 2026. [Chart deception in vision-language models: From vulnerability to mitigation](#). *Preprint*, arXiv:2607.22600.
- Khandaker Abrar Nadib, Marina Kogan, Alexander Lex, and Maxim Lisnic. 2026. [Guardrail selection in line charts to contextualize persuasive visualizations](#). In *Computer Graphics Forum (EuroVis)*, 45(3), 2026.
- Swaroop Panda. 2026. [Making ai agents evaluate misleading charts without nudging](#).
- Christian Seto, Jacqueline Nguyen, Jiayi Hong, and Ross Maciejewski. 2026. [Llms have visualization literacy: Now what? experiments exploring llm visualization evaluation capabilities](#). In *arXiv (Cornell University)*. Cornell University.
- Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. [Reflexion: Language agents with verbal reinforcement learning](#). *Preprint*, arXiv:2303.11366.
- Jonathan Tonglet, Tinne Tuytelaars, Marie-Francine Moens, and Iryna Gurevych. 2026. [Protecting multimodal large language models against misleading visualizations](#). In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8329–8349, San Diego, California, United States. Association for Computational Linguistics.
- Jonathan Tonglet, Jan Zimny, Tinne Tuytelaars, and Iryna Gurevych. 2025. [Is this chart lying to me? automating the detection of misleading visualizations](#). arXiv:2508.21675. *Preprint*, arXiv:2508.21675.
- Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yuan Yang, Jingsen Zhang, Zhiyuan Chen, Jikai Tang, Xu Chen, Yankai Lin, Wayne Xin Zhao, Zhewei Wei, and Ji-Rong Wen. 2024. [A survey on large language model based autonomous agents](#). *Frontiers of Computer Science*, 18(6).
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2022. [Self-consistency improves chain of thought reasoning in language models](#).
- Zhun Wang, Vincent Siu, Zhe Ye, Tianneng Shi, Yuzhou Nie, Xuandong Zhao, Chenguang Wang, Wenbo Guo, and Dawn Song. 2025. [Agentvigil: Automatic black-box red-teaming for indirect prompt injection against llm agents](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 23159–23172, Suzhou, China. Association for Computational Linguistics.
- Chen Henry Wu, Rishi Shah, Jing Yu Koh, Ruslan Salakhutdinov, Daniel Fried, and Aditi Raghunathan. 2024. [Dissecting adversarial robustness of multimodal llm agents](#). *Preprint*, arXiv:2406.12814.
- Qiusi Zhan, Zhixiang Liang, Zifan Ying, and Daniel Kang. 2024. [Injecagent: Benchmarking indirect prompt injections in tool-integrated large language model agents](#). In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 10471–10506, Bangkok, Thailand. Association for Computational Linguistics.
- Yanjie Zhang, Yafei Li, Rui Sheng, Zixin Chen, Yanna Lin, Huamin Qu, Lei Chen, and Yushi Sun. 2026. [Navigating the mirage: A dual-path agentic framework for robust misleading chart question answering](#). In *arXiv (Cornell University)*. Cornell University.
- Wei Zhao, Zhe Li, Peixin Zhang, and Jun Sun. 2026. [Clawguard: A runtime security framework for tool-augmented llm agents against indirect prompt injection](#). *Preprint*, arXiv:2604.11790.

A Execution Loop and Worked Example

Algorithm 1 summarizes one execution of CERS under the candidate budget K and the review budget L for a single verification node.

A.1 A Worked Example of Revision and Cross-Step Reuse

We illustrate the mechanism on the chart semantics of task pub013 from our execution data, which asks the agent to choose the highest-risk record among Illinois, Kansas, and Delaware; the case material records that the map legend assigns lighter yellow to higher risk and deep red to lower risk. The walkthrough below demonstrates the method and does not add experimental results. Suppose the current proposal selects Kansas under

$$\pi_0 : \underbrace{\text{Kansas is darker}}_{O_0} \wedge \underbrace{\text{darker means higher risk}}_{B_0} \Rightarrow C_q(\text{KS}). \quad (9)$$

The competing module asks why a darker color should mean a higher risk; the model generates legend-dependent interpretations on the same complete map and compares the legend with the default color convention. If verification confirms that the chart does not support B_0 , the rule loses applicability on this map and the state records the original interpretation, the revision reason that the convention contradicts the current legend definition, the current interpretation that candidate colors are read through the verified legend mapping, the scope of the current map and risk metric, and the condition

for further updates, namely a different encoding or an effective refutation of the legend reading. The agent then re-compares the candidates with B_1 and the observed colors and completes the judgment the public task requires; the concrete option is an output of this reasoning, while the rule state keeps the interpretation basis of the current map. When the agent enters later pages or re-revisits the chart-dependent choice, what it reads is that rule and its revision reason: a new explanation that merely repeats “deep red usually means more danger” must answer the recorded legend evidence and cannot restore B_0 directly, while a re-reading that finds the earlier legend verification mistaken is allowed to revise again. The process records three distinct things: what changed in the interpretation rules, whether the current action must change, and which rule version later steps actually used.

B Additional Results and Boundary Statements

B.1 Development Subset and Stratified Results

The 24-task development subset is scored separately in every round. In the original round the full configuration V5 gains +17 paired decisions on the remaining 116 tasks and +1 on the development subset; the same decomposition holds for the 8B round (+10 splits as +9 and +1). The environment domain is the only one in which V5 loses against direct selection in both model rounds (−5 in the 27B replication, −9 in the 8B round). At the family level, subset-extrapolation tasks contribute +18 (original round) and +11 (8B round) for V5; dual-axis and non-linear-axis families are mildly negative for V5 in the second model round (−5 and −6), and bar-height tasks are near-neutral (−3 and −2).

B.2 Interface-Failure Accounting

Interface failures are runs in which the agent never produces a schema-valid selection. They are zero for Direct, V4, V5, V6, and V7 on the 27B model (all three rounds), and 6, 8, and 51 for V3 in the original, replication, and 8B rounds respectively. V3’s interface failures are not discarded: they remain in the denominator, which is why its large error-to-target correction counts on the 27B model (30 forward against 6 reverse, exact McNemar $p = 0.0001$) reduce to a net +2. The 8B actor’s failure mode is dominated by responses that violate

the typed-record contract of the verification pass rather than by timeouts or truncations.

B.3 Online Deployment Details

Each online unit executes one task on one chart variant under one system inside a sealed environment: public inputs are mounted read-only, rule hooks intercept chart evidence entering the context, and the unit ends in exactly one of the recorded terminal statuses (submitted, abstention with reason, re-verification budget exhausted, or engineering failure). The strict interface (`typed_scalars.v2`) permits at most eight actor calls and one re-verification round per unit at temperature zero. The three actor systems differ only below the interface: a local 27B model under strict response formats (Batch A), the same model behind a schema-guided proxy mirroring the sealed structural rules (Batch B), and the frontier model reached directly through the gateway (Batch C and its stability rerun). Across the two frontier-model batches the defense arm read its persistent rule state 48 and 46 times across its 24 units, never cooccurred with a wrong submission, and every per-unit persistence-causality flag reads `not_established`; the ordinary arm produced exactly six wrong submissions, all chart-induced, all under the 27B actor on misleading charts, and none under the frontier actor.

B.4 Boundary Statements

Three boundary statements delimit the claims. First, the per-unit `causal_persistence` effect is recorded as *not established* in every online unit: the harness shows that revised rules were re-read and that outcomes changed between arms, but a single-unit causal attribution of outcome change to persistence is not identified by this design. Second, the stability rerun’s unit-level agreement (37/48) bounds the reliability of any single-batch online estimate; disaggregated four-way cells shift by at most two units, and no flip ever crosses into a wrong submission. Third, the online deployment’s paired effect is one-directional: the defense arm converts ordinary decisions into abstentions (10 and 13 per batch) and never the reverse, so the deployment should be read as a safe-fail characterization of the defense under the strict protocol, not as an accuracy improvement result.

Algorithm 1 CERS execution loop

- 1: inputs: task q ; base agent $\mathcal{F}_\theta$ and the original action interface; chain budget K ; review budget L .
 - 2: initialize rule state $M \leftarrow \emptyset$; execution history $H \leftarrow \emptyset$.
 - 3: **for** each step t **do**
 - 4: receive observation o_t ; update chart references and interpretation context c_t .
 - 5: read applicable rules, revision evidence, and disputes from M .
 - 6: $(\tilde{a}_t, \tilde{\pi}_t) \leftarrow \mathcal{F}_\theta(q, o_t, H, \mathcal{R}_t)$.
 - 7: **if** a trigger holds: first chart-dependent judgment, interpretation conflict, or applicability change **then**
 - 8: expand the current explanation into $O \wedge B \Rightarrow C_q(\tilde{a}_t)$.
 - 9: generate at most K substantially different competing chains on the same complete chart.
 - 10: normalize candidates; merge repeated chains and shared rules.
 - 11: verify observations, rule applicability, and derivation sufficiency.
 - 12: apply adoption, revocation, replacement, or scope adjustment by the evidence.
 - 13: re-derive the current action from the updated state; within L , resolve remaining key disputes.
 - 14: **end if**
 - 15: if the proposal reuses a revoked rule without new evidence, return the revision record and require re-derivation; with an effective new challenge, re-verify the rule.
 - 16: with KEEP or REVISE, execute the approved action through the original interface; with exhausted budget, record UNRESOLVED and stop the decision.
 - 17: append the true observation, proposal, rule versions, action, and feedback to H .
 - 18: **end for**
 - 19: save the task’s rule revisions and execution trajectory; close the task’s rule scope.
-