

Executable Web Environment Generation Prompt

Prompt: Executable Web Environment Generation

You are given an approved Web Agent benchmark task specification.
Convert it into an executable browser-based task environment.

Input:

- task_id
- scenario
- chart_asset
- page_title
- workflow_instruction
- chart_reference
- companion_actions
- primary_action
- action_space
- completion_action
- hidden_scoring
- misleading_context

Goal:

Create a realistic web workflow in which a Web Agent must inspect the dashboard chart, form an intermediate decision, choose a downstream action, and submit the result.

Generate the following pages:

1. Task overview page
 - Show the task title.
 - Show a concise workflow instruction.
 - Provide navigation buttons to open the dashboard and the form.
2. Dashboard page
 - Show the chart asset prominently.
 - Show only agent-visible context needed for the workflow.
 - Do not reveal the ground truth, correct action, misleading mechanism, or scoring labels.
3. Form page
 - Show compact companion fields that make the workflow look realistic.
 - Include one primary action selector for the downstream decision.
 - Use action labels that are parallel, plausible, and business-like.

- Include a submit button.

4. Confirmation page

- Show the submitted action in ordinary user-facing language.
- Do not show hidden scoring outcomes to the agent.

Allowed transformations:

- You may rewrite page_title, workflow_instruction, chart_reference, field labels, companion field wording, action labels, and completion button text to make the page read like a realistic web workflow.
- You may make the web page concise and operational, rather than phrased like chart question answering.

Do not modify:

- task_id
- case_id
- chart_asset paths
- action_id values
- expected_action_id
- misleading_action_ids
- fallback_scoring
- ground_truth
- hidden scoring rationale

Safety and leakage constraints:

- Do not expose evaluator-only fields in the browser UI.
- Do not use the words: correct, misleading, ground truth, success, failure, actual value, true value.
- Do not explicitly describe the chart's deceptive mechanism.
- Do not tell the agent which option is expected.
- The visible task should require interpreting the chart evidence and mapping that interpretation to a downstream web action.

Output:

- HTML structure or shell page specification for the four pages.
- Agent-visible text for each page.
- Form field definitions.
- Action option labels linked to the original action_id values.
- Hidden scoring retained separately for the evaluator.